



DREXEL UNIVERSITY

Electrical and Computer Engineering

College of Engineering

Drexel University

Electrical and Computer Engineering Dept.

Parallel Computer Architecture ECEC-622

TITLE: SSE Assignment: Gaussian Elimination

GROUP MEMBERS: Gregory Matthews and Mark Klobukov

INSTRUCTOR: Dr. Kandasamy

DATE SUBMITTED: 2/28/2017

DATE DUE: 2/28/2017

Assignment description:

The first PThread assignment involved parallelizing a given program that performs Gaussian elimination of an $N \times N$ system of linear equations. The algorithm is shown in the pseudo-code below:

```
1: procedure GAUSS_ELIMINATE( $A, b, y$ )
2:   int  $i, j, k$ ;
3:   for  $k := 0$  to  $n - 1$  do
4:     for  $j := k + 1$  to  $n - 1$  do
5:        $A[k, j] := A[k, j] / A[k, k]$ ;    /* Division step. */
6:     end for
7:      $y[k] := b[k] / A[k, k]$ ;
8:      $A[k, k] := 1$ ;
9:     for  $i := k + 1$  to  $n - 1$  do
10:      for  $j := k + 1$  to  $n - 1$  do
11:         $A[i, j] := A[i, j] - A[i, k] \times A[k, j]$ ;    /* Elimination step. */
12:      end for
13:       $b[i] := b[i] - A[i, k] \times y[k]$ ;
14:       $A[i, k] := 0$ ;
15:    end for
16:  end for
```

Parallelization approach:

To begin the SSE Assignment, we had to consider that the address locations for allocating space on the virtual address space has to be 16 byte aligned, as per SSE's requirements when working with a 128 bit XMM registers and attempting to store 4 floats which are 4 bytes each, or 32 bits. 4 floats stored into an XMM register will allow for the simultaneous operation of 4 floats in parallel. By creating 16 byte aligned elements in each Matrix structure for Gaussian Elimination, parallelization of each row in the division and elimination steps can be parallelized by storing 4 row elements each iteration of the j loop, and performing the necessary division and elimination operations on those stored floats held in each XMM registers, and loaded back into the Matrix upon successful operation.

Compile with `gcc -o gauss_eliminate gauss_eliminate.c compute_gold.c -std=c99 -O3 -lm`

SSE code explanation:

Division and elimination step were both included in the **gauss_eliminate_using_sse()** function, the screenshots of which are included and described below.

```
96 void
97 gauss_eliminate_using_sse(float* U, const float* A, unsigned int num_elements)
98 {
99     unsigned int i, j, k, init;
100
101     __m128 m0, m1, m2, m3;
102     // INITIALIZATION
103     for (i = 0; i < num_elements; i++) /* Copy the contents of the A matrix into t
104         for (j = 0; j < num_elements/ 4; j++){
105             m0 = _mm_load_ps(&A[num_elements*i + 4*j]);
106             _mm_store_ps(&U[num_elements*i + 4*j], m0);
107         }
```

The above screenshot shows that the function accepts pointers to the elements of matrices U and A, as well as the number of elements. The argument types are the same as in the **compute_gold()** function provided by Dr. Kandasamy. On line 101, four **__m128** variables are declared. Then, on line 103, the contents of matrix A are copied into matrix U. The matrices were created such that they are byte aligned (using **posix_memalign()** function), and since the number of elements in each row is divisible by 4, it is guaranteed that we can access every 4th element and utilize SSE registers to speed up the copying process.

```
110     for (k = 0; k < num_elements; k++){
111
112         m0 = _mm_set_ps1(U[num_elements*k + k]);
113         init = k+1;
114
115         // INITIAL SERIAL DIVISION
116         while ((init % 4) != 0){
117             U[num_elements * k + init] = (U[num_elements * k + init] / U[num_elements * k + k]);
118             init++;
119         }
120
121         // SSE DIVISION STEP
122         for (j = init; j < num_elements; j+=4){ /* Reduce the current row. */
123
124             m1 = _mm_load_ps(&U[num_elements*k + j]);
125             m1 = _mm_div_ps(m1,m0);
126             _mm_store_ps(&U[num_elements*k + j], m1);
127         }
128
129         U[num_elements * k + k] = 1; /* Set the principal diagonal entry in U to be 1. */
130     }
```

On line 110, the outermost loop with the counter **k** (like in the pseudo-code algorithm description) is started. In the division step, each element of a given row is divided by the corresponding pivot element. Divisions are performed in groups of 4 (128-bit SSE registers fit 4 floats at a time), so the pivot element needs to be loaded into all 4 words of **m0**. This is achieved with the function **_mm_set_ps1()**.

It was necessary to make sure that the number of elements in a given row that are divided using SSE is divisible by 4. For this reason, the while-loop on line 116 was included. It performs the division step serially until an element whose index is divisible by 4 is reached. This index is stored in the **init** variable. The rest of the division step is done with SSE in the for-loop starting on line 122.

There are three statements in the for loop. The first one on line 124 loads four floats from the current matrix row into **m1**. The next line divides these four floats by the pivot element of the given row. The last command in the for loop stores the division back into the result matrix U.

Finally, the principal diagonal entry in U is set to 1, as on line 8 of the pseudocode algorithm.

```

132 // ELIMINATION STEP
133 for (i = (k+1); i < num_elements; i++){
134
135     init = k+1;
136
137     // INITIAL SERIAL ELIMINATION
138     while ((init % 4) != 0){
139         U[num_elements * i + init] = U[num_elements * i + init] \-
140             (U[num_elements * i + k] * U[num_elements * k + init]);
141         init++;
142     }
143
144     m0 = _mm_load1_ps(&U[num_elements*i + k]);
145
146     // SSE ELIMINATION
147     for (j = init; j < num_elements; j+=4){
148
149         m2 = _mm_load_ps(&U[num_elements*i + j]);
150         m3 = _mm_load_ps(&U[num_elements*k + j]);
151         m3 = _mm_mul_ps(m3, m0);
152         m2 = _mm_sub_ps(m2, m3);
153         _mm_store_ps(&U[num_elements*i + j], m2);
154     }
155
156     U[num_elements * i + k] = 0;
157 }
158 }
159 }
160 }
161

```

Much the same ideas that were used in parallelizing the division step, the elimination step, shown in the above screenshot, also had a conditional while loop statement on line 146, which checks that the current row index is divisible by 4, and if so, continue the elimination

execution serially. This index is not parallelizable, and we must wait until the starting row index is byte aligned correctly by checking until it is divisible by 4, and then continue SSE execution using the variable `init`, which will now be divisible by 4 when it exits the while loop. At the beginning of each `i` iteration, we load `U[num_elements*i + k]` by using `_mm_load_ps` into a 128 bit XMM register for use later inside the `j` loop. Inside the `j` loop we also load `U[num_elements*i + k]` and `U[num_elements*k+j]` into registers to perform the multiplication and subtraction operations necessary for the elimination algorithm. Once fully computed and stored into an XMM register, this value is stored back into the `U` Matrix at address `U[num_elements*i + j]`.

Modification of the `allocate_matrix()` function:

In order to use SSE for matrix elimination, the `allocate_matrix()` function was changed such that the memory allocated to the matrices is 16 byte aligned. Instead of `malloc()`, `posix_memalign()` was used to allocate memory. This is shown on lines 193-196 of the code snippet below.

```
186 Matrix
187 allocate_matrix(int num_rows, int num_columns, int init){
188     Matrix M;
189     M.num_columns = M.pitch = num_columns;
190     M.num_rows = num_rows;
191     int size = M.num_rows * M.num_columns;
192
193     void* allocation;
194     posix_memalign(&allocation, 16, sizeof(float) * size);
195
196     M.elements = (float*) allocation;
197
198     for(unsigned int i = 0; i < size; i++){
199         if(init == 0) M.elements[i] = 0;
200         else
201             M.elements[i] = get_random_number(MIN_NUMBER, MAX_NUMBER);
202     }
203
204     return M;
205 }
```

Result:

The following table shows serial time, SSE run time, and the speedup for 4 different matrix sizes. Although the assignment instructions specified that the matrix is guaranteed to be 2048x2048 elements, we made our code more flexible so that it can handle other matrix sizes as long as they are divisible by 4.

Table 1: Run Time Comparison between Serial and SSE versions of the program

Matrix size	Time Serial (s)	Time SSE (s)	Speedup
512x512	0.09	0.02	4.50
1024x1024	0.57	0.16	3.56
2048x2048	3.76	1.3	2.89
4096x4096	31.86	13.05	2.44




Figure 1: Speedup vs Matrix Size